



UNIVERSITY OF MISSOURI-KANSAS CITY

Bigdata Hadoop Programming

Lab 1 Assignment

Team Members:

Pragathi Thammaneni

Sridevi Mallipudi

Introduction:

The main core concept for executing the Lab 1 Assignment is to implement the Map Reduce Algorithm for finding the solution for the Facebook common friends problem and to run the MapReduce job on Apache Hadoop. And implementing the No Sql comparison over HBase and Cassandra.

Objectives:

To code for the 2 questions the below concepts are implemented.

Map Reduce Algorithm for Facebook common friend's problem

Cassandra

HBase

No SQL Comparison over HBase and Cassandra

Approaches /Methods:

Using Hadoop 1.2.1, HBase 0.94.8, IntelliJ (Community edition)

Workflow &Datasets/Parameters and Evaluation:

The below each question will follow different approaches to solve. Coding is done to perform the evaluation of each individual snippet to execute the datasets which are provided as the input parameters.

Question 1:

1. Hadoop MapReduce Algorithm

Implement MapReduce algorithm for finding Facebook common friends problem and run the MapReduce job on Apache Hadoop. Show your implementation through map-reduce diagram as shown in Lesson Plan 2 : (<https://umkc.box.com/s/epp04s3m6g8jw6ulnnb4meqgnbwzbz5uc>) Write a report including your algorithm and result screenshots.

Finding Facebook common friends: Facebook has a list of friends (note that friends are a bi-directional thing on Facebook. If I'm your friend, you're mine). They also have lots of disk space and they serve hundreds of millions of requests everyday. They've decided to pre-compute calculations when they can to reduce the processing time of requests. One common processing request is the "You and Joe have 230 friends in common" feature. When you visit someone's profile, you see a list of friends that you have in common. We're going to use MapReduce so that we can calculate everyone's common friends once a day and store those results. Later on it's just a quick lookup. We've got lots of disk, it's cheap.

Example (What is the Key/Value Pair?)

Assume the friends are stored as Person -> [List of Friends], our friends list is then:

A -> B C D
B -> A C D E
C -> A B D E
D -> A B C E
E -> B C D

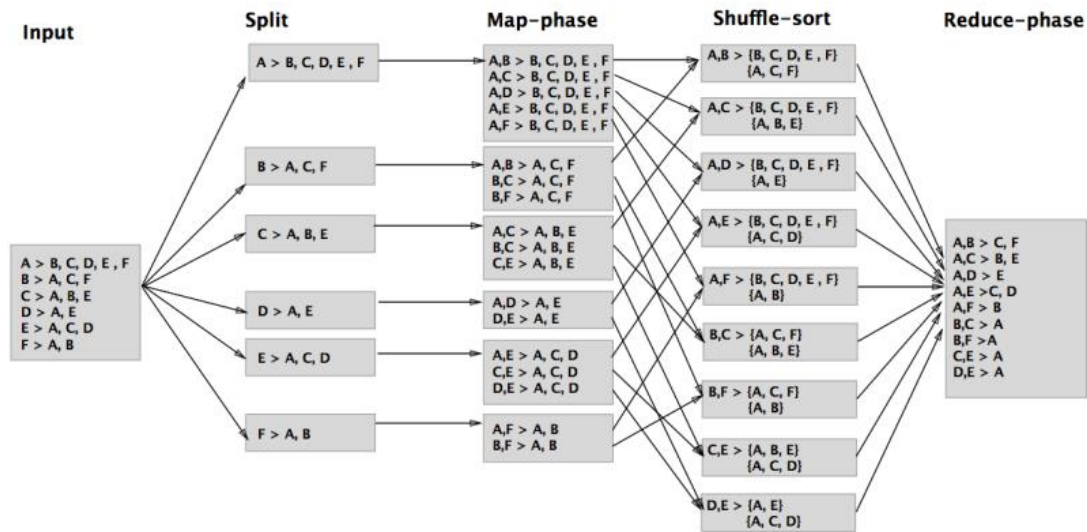
The result after reduction is:

(A B) -> (C D)
(A C) -> (B D)
(A D) -> (B C)
(B C) -> (A D E)
(B D) -> (A C E)
(B E) -> (C D)
(C D) -> (A B E)
(C E) -> (B D)
(D E) -> (B C)

When D visits B's profile, we can quickly look up (B D) and see that they have three friends in common, (A C E).

For solving the below question a map reduce algorithm is used to find out the face book common friend problem. Below is the sample use case for implementing the Map reduce .

Use Case Diagram:

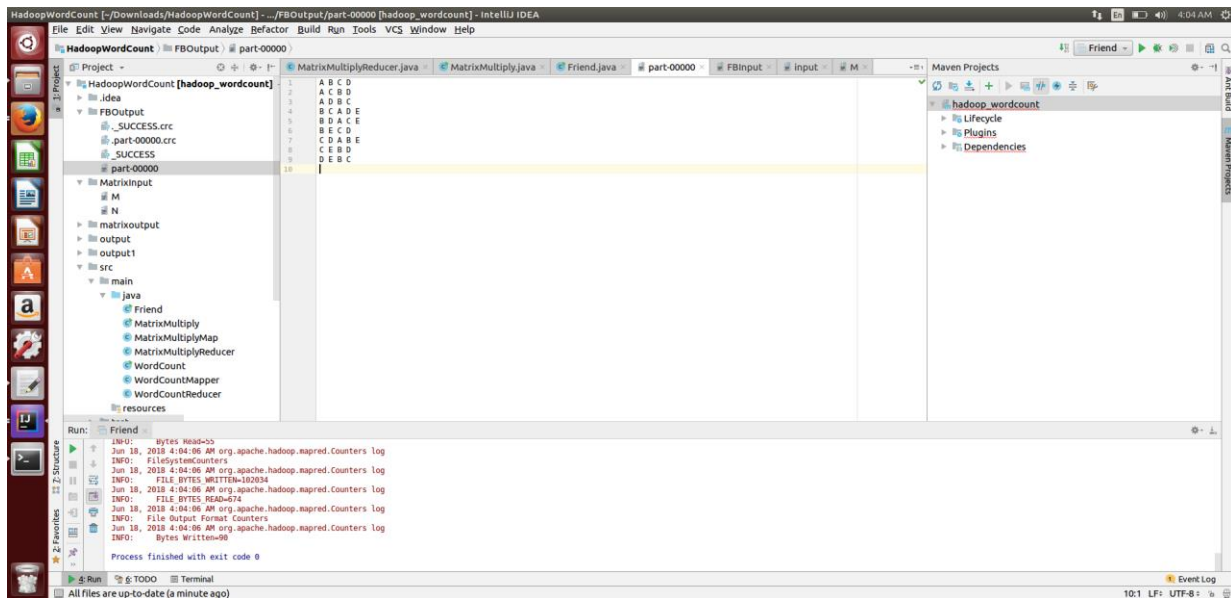


Below is the screen shot for the Map Reduce Algorithm

The screenshot shows an IDE with the following components:

- Project Explorer:** Shows the project structure with folders like 'src', 'resources', and 'test'. The 'src' folder contains the 'main' package with classes like 'Friend', 'MatrixMultiply', 'MatrixMultiplyMap', 'WordCount', 'WordCountMapper', and 'WordCountReducer'.
- Code Editor:** Displays the 'Friend.java' file. It shows imports for various Java and Hadoop classes, followed by the implementation of the 'Friend' class, which implements the 'Mapper' interface. The 'map' method processes the input text and writes it to the output collector.
- Run Console:** Shows the output of the program execution, including logs for 'Bytes Read', 'FilesystemCounters', 'File Bytes Written', 'File Bytes Read', 'File Output Format Counters', and 'Bytes Written'.

Output Screenshot:



Question 2:

2. Use Case Based No SQL Comparison

Consider one of the use cases from the below link:

<https://umkc.box.com/s/q64fvjm6yvd454w5v3ky0he4854g6m1fq>

These use cases were discussed in Lecture 1: Cassandra.

- Consider one of the use case and use a simple dataset. Describe the use case considered based on your assumptions, report the dataset, its fields, datatype etc.
- Use HBase to implement a Solution for the use case. Report at least 3 queries, their input and output. The query's relevance towards solving the use case is important.
- Use Cassandra to implement a Solution for the use case. Report at least 3 queries, their input and output. The query's relevance towards solving the use case is important.
- Compare Cassandra and HBase for your use case. Present a table with comparison of your use case being implemented in both NO SQL Systems.

For solving the above 4 cases below work flow is been followed, created tables in Hbase and Cassandra and then performed three queries on each. Also Performed the comparison over Cassandra and Hbase for the usecase being implemented in both No SQL Systems.

No Sql Comparison over Hbase and Cassandra

CASSANDRA implementation:

For solving this use case we opt to choose Coursera data set, the following are the fields and data types used for implementing the Cassandra.

The data set will include below fields and data types

```
create keyspace coursera with replication={'class':'SimpleStrategy','replication_factor':2};
cqlsh> use coursera;
```

```
// Creating the table with required
create table if not exists coursera_details (
  learner_id int PRIMARY KEY,
  course_id int,
  course_name text,
  course_duartion int,
  tutor_name text);
```

Created a table with the above details so that the below CRUD operations can be implemented to retrieve the data

Update ,Select and delete operations are performed on the data set .Below are the commands and respective screen shorts for the execution

Queries for implementing the Cassandra:

```
-----CASSANDRA Use case -----
create keyspace coursera with replication={'class':'SimpleStrategy','replication_factor':2};
cqlsh> use coursera;

// Creating the table with required
create table if not exists coursera_details (
  learner_id int PRIMARY KEY,
  course_id int,
  course_name text,
  course_duartion int,
  tutor_name text);

// Insert commands
INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (1,123,'Bigdata',6,'Jack');
INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (2,124,'Software Engineering',7,'Jim');

INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (3,125,'Statistical Learning',2,'Sam');

INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (4,126,'Cloud Computing',6,'Bill');

INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (5,127,'Hadoop eco system',6,'Roy');
INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values (6,128,'Cloud Computing',4,'Lee');

Query - 1 (UPDATE DATA)
//Update statement for updating the course details FROM course_duartion ~2 ,course_name Statistical Learning TO 5 and Introduction to Statistical Learning
Update coursera.coursera_details set course_duartion =5 ,course_name='Introduction to Statistical Learning' where learner_id=3;

Query - 2 (SELECT DATA)
// Select statement for reading the data with filtering
select tutor_name from coursera.coursera_details where course_name='Cloud Computing' ALLOW FILTERING;

Query - 3 (DELETE DATA)
// Deleting the data from the table
DELETE from coursera.coursera_details where learner_id = 2;
```

Execution for Cassandra use case:

Table Creation

```
Cassandra CQL Shell

WARNING: console codepage must be set to cp65001 to support utf-8 encoding on Windows platforms.
If you experience encoding problems, change your console codepage with 'chcp 65001' before starting cqlsh.

Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.0.9 | CQL spec 3.4.0 | Native protocol v4]
Use HELP for help.
WARNING: pyreadline dependency missing. Install to enable tab completion.
cqlsh> create keyspace coursera with replication={'class': 'SimpleStrategy', 'replication_factor': 2};
cqlsh> use coursera;
cqlsh:coursera> create table if not exists coursera_details (
    ... learner_id int PRIMARY KEY,
    ... course_id int,
    ... course_name text,
    ... course_duartion int,
    ... tutor_name text);
cqlsh:coursera>
```

Insert Statements

```
Cassandra CQL Shell

cqlsh:coursera>
cqlsh:coursera> INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values
(3,125,'Statstical Learning',2,'Sam');
cqlsh:coursera>
cqlsh:coursera> INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values
(4,126,'Cloud Computing',6,'Bill');
cqlsh:coursera>
cqlsh:coursera> INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values
(5,127,'Hadoop eco system',6,'Roy');
cqlsh:coursera>
cqlsh:coursera> INSERT into coursera.coursera_details(learner_id,course_id,course_name,course_duartion,tutor_name)values
(6,128,'Cloud Computing',4,'Lee');
cqlsh:coursera> select * from coursera.coursera_details;

 learner_id | course_duartion | course_id | course_name          | tutor_name
-----|-----|-----|-----|-----
          5 |                6 |        127 | Hadoop eco system   | Roy
          1 |                6 |        123 | Bigdata             | Jack
          2 |                7 |        124 | Software Engineering | Jim
          4 |                6 |        126 | Cloud Computing     | Bill
          6 |                4 |        128 | Cloud Computing     | Lee
          3 |                2 |        125 | Statstical Learning | Sam

(6 rows)
cqlsh:coursera>
```

Query 1- UPDATE DATA

```
Cassandra CQL Shell
cqlsh:coursea>
cqlsh:coursea> INSERT into coursea.coursea_details(learner_id,course_id,course_name,course_duartion,tutor_name)values
(6,128,'Cloud Computing',4,'Lee');
cqlsh:coursea> select * from coursea.coursea_details;

learner_id | course_duartion | course_id | course_name          | tutor_name
-----+-----+-----+-----+-----
5 | 6 | 127 | Hadoop eco system | Roy
1 | 6 | 123 | Bigdata | Jack
2 | 7 | 124 | Software Engineering | Jim
4 | 6 | 126 | Cloud Computing | Bill
6 | 4 | 128 | Cloud Computing | Lee
3 | 2 | 125 | Statstical Learning | Sam

(6 rows)
cqlsh:coursea> Update coursea.coursea_details set course_duartion =5 ,course_name='Introduction to Statstical Learning' where learner_id=3;
cqlsh:coursea> select * from coursea.coursea_details;

learner_id | course_duartion | course_id | course_name          | tutor_name
-----+-----+-----+-----+-----
5 | 6 | 127 | Hadoop eco system | Roy
1 | 6 | 123 | Bigdata | Jack
2 | 7 | 124 | Software Engineering | Jim
4 | 6 | 126 | Cloud Computing | Bill
6 | 4 | 128 | Cloud Computing | Lee
3 | 5 | 125 | Introduction to Statstical Learning | Sam

(6 rows)
cqlsh:coursea>
```

Query 2- SELECT DATA

```
Cassandra CQL Shell
(6 rows)
cqlsh:coursea> select tutor_name from coursea.coursea_details where course_name='Cloud Computing' ALLOW FILTERING;

tutor_name
-----
Bill
Lee

(2 rows)
cqlsh:coursea>
```

Query 3- DELETE DATA

```
Cassandra CQL Shell
cqlsh:courseera> DELETE from courseera.courseera_details where learner_id = 2;
cqlsh:courseera> select * from courseera.courseera_details;
```

learner_id	course_duartion	course_id	course_name	tutor_name
5	6	127	Hadoop eco system	Roy
1	6	123	Bigdata	Jack
4	6	126	Cloud Computing	Bill
6	4	128	Cloud Computing	Lee
3	5	125	Introduction to Statstical Learning	Sam

```
(5 rows)
cqlsh:courseera>
```

HBASE Implementation:

For solving this use case we opt to choose Couresera data set, the following are the fields and data types used for implementing the Hbase.


```
create 'coursera', 'coursera_details'

> put 'coursera', '1', 'coursera_details:learner_Id', '1'
> put 'coursera', '1', 'coursera_details:course_id', '123'
> put 'coursera', '1', 'coursera_details:course_name', 'Bigdata'
> put 'coursera', '1', 'coursera_details:course_duration', '6'
> put 'coursera', '1', 'coursera_details:tutor_name', 'Jack'
> put 'coursera', '2', 'coursera_details:learner_Id', '2'
> put 'coursera', '2', 'coursera_details:course_id', '124'
> put 'coursera', '2', 'coursera_details:course_name', 'Software Engineering'
> put 'coursera', '2', 'coursera_details:course_duration', '7'
> put 'coursera', '2', 'coursera_details:tutor_name', 'Jim'
```

updating

```
put 'coursera', '2', 'coursera_details:course_name', 'Sattistical Learning'
```

Selecting:

```
get 'coursera', '1'
get 'coursera', '2'
```

Deleting:

```
deleteall 'coursera', '1'
```

Output Screen shots for HBase:

Table creation:

```
sridevi@sridevi-VirtualBox:~$ hbase(main):005:0> put 'coursera', '1', 'coursera_details:tutor_name', 'Jack'
0 row(s) in 0.0170 seconds
hbase(main):005:0> put 'coursera', '1', 'coursera_details:course_id', '123'
0 row(s) in 0.0260 seconds
hbase(main):006:0> put 'coursera', '1', 'coursera_details:course_name', 'Bigdata'
0 row(s) in 0.0250 seconds
hbase(main):007:0> put 'coursera', '1', 'coursera_details:course_duration', '6'
0 row(s) in 0.0150 seconds
hbase(main):008:0> put 'coursera', '2', 'coursera_details:learner_Id', '2'
0 row(s) in 0.0250 seconds
hbase(main):009:0> put 'coursera', '2', 'coursera_details:course_id', '124'
0 row(s) in 0.0150 seconds
hbase(main):010:0> put 'coursera', '2', 'coursera_details:course_name', 'Software Engineering'
0 row(s) in 0.0250 seconds
hbase(main):011:0> put 'coursera', '2', 'coursera_details:course_duration', '7'
0 row(s) in 0.0490 seconds
hbase(main):012:0> put 'coursera', '2', 'coursera_details:tutor_name', 'Jim'
hbase(main):013:0> describe 'coursera'
hbase(main):014:0> list
hbase(main):015:0> sridevi@sridevi-VirtualBox:~$ hbase shell
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.94.8, r1485467, Wed May 22 20:53:13 UTC 2013
hbase(main):001:0> scan 'coursera'
ROW
1
1
1
1
2
2
2
2
2 row(s) in 1.7920 seconds
hbase(main):002:0> describe 'coursera'
DESCRIPTION
ENABLED
'coursera', {NAME => 'coursera_details', ENCODE_ON_DISK => 'true', BLOOMFILTER => 'NONE', VERSIONS => '3', IN_MEMORY => 'false', KE true
EP_DELETED_CELLS => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => '2147483647', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKC
ACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.7250 seconds
hbase(main):003:0>
```

Query 1- UPDATE DATA

```
srdevi@srdevi-VirtualBox:~$ hbase(main):010:0> put 'coursera', '2', 'coursera_details:tutor_name', 'JlIn
hbase(main):011:0> scan 'coursera'
hbase(main):012:0> describe 'coursera'
hbase(main):013:0>
hbase(main):014:0> list
hbase(main):015:0> srdevi@srdevi-VirtualBox:~$ hbase shell
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.94.0, r1485407, Wed May 22 20:53:13 UTC 2013

hbase(main):001:0> scan 'coursera'
ROW
1      column=coursera_details:course_duration, timestamp=1529314663277, value=6
1      column=coursera_details:course_id, timestamp=1529314641749, value=123
1      column=coursera_details:course_name, timestamp=1529314651864, value=Bigdata
1      column=coursera_details:learner_id, timestamp=1529314629019, value=1
1      column=coursera_details:tutor_name, timestamp=1529314672629, value=Jack
2      column=coursera_details:course_duration, timestamp=1529314738104, value=7
2      column=coursera_details:course_id, timestamp=1529314708322, value=124
2      column=coursera_details:course_name, timestamp=1529314725389, value=Software Engineering
2      column=coursera_details:learner_id, timestamp=1529314684719, value=2
2 row(s) in 1.7920 seconds

hbase(main):002:0> describe 'coursera'
DESCRIPTION
'coursera', {NAME => 'coursera_details', ENCODE_ON_DISK => 'true', BLOOMFILTER => 'NONE', VERSIONS => '3', IN_MEMORY => 'false', KE true
EP_DELETED_CELLS => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => '2147483647', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKC
ACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.7250 seconds

hbase(main):003:0> put 'coursera', '2', 'coursera_details:course_name', 'Sattist
ical Learning'
0 row(s) in 0.1570 seconds

hbase(main):004:0> scan 'coursera'
ROW
1      column=coursera_details:course_duration, timestamp=1529314663277, value=6
1      column=coursera_details:course_id, timestamp=1529314641749, value=123
1      column=coursera_details:course_name, timestamp=1529314651864, value=Bigdata
1      column=coursera_details:learner_id, timestamp=1529314629019, value=1
1      column=coursera_details:tutor_name, timestamp=1529314672629, value=Jack
2      column=coursera_details:course_duration, timestamp=1529314738104, value=7
2      column=coursera_details:course_id, timestamp=1529314708322, value=124
2      column=coursera_details:course_name, timestamp=1529315359923, value=Sattistical Learning
2      column=coursera_details:learner_id, timestamp=1529314684719, value=2
2 row(s) in 0.1430 seconds

hbase(main):005:0>
```

Query 2- SELECT DATA

```
srdevi@srdevi-VirtualBox:~$ hbase(main):001:0> scan 'coursera'
ROW
2      column=coursera_details:learner_id, timestamp=1529314684719, value=2
2 row(s) in 1.7920 seconds

hbase(main):002:0> describe 'coursera'
DESCRIPTION
'coursera', {NAME => 'coursera_details', ENCODE_ON_DISK => 'true', BLOOMFILTER => 'NONE', VERSIONS => '3', IN_MEMORY => 'false', KE true
EP_DELETED_CELLS => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => '2147483647', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKC
ACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.7250 seconds

hbase(main):003:0> put 'coursera', '2', 'coursera_details:course_name', 'Sattist
ical Learning'
0 row(s) in 0.1570 seconds

hbase(main):004:0> scan 'coursera'
ROW
1      column=coursera_details:course_duration, timestamp=1529314663277, value=6
1      column=coursera_details:course_id, timestamp=1529314641749, value=123
1      column=coursera_details:course_name, timestamp=1529314651864, value=Bigdata
1      column=coursera_details:learner_id, timestamp=1529314629019, value=1
1      column=coursera_details:tutor_name, timestamp=1529314672629, value=Jack
2      column=coursera_details:course_duration, timestamp=1529314738104, value=7
2      column=coursera_details:course_id, timestamp=1529314708322, value=124
2      column=coursera_details:course_name, timestamp=1529315359923, value=Sattistical Learning
2      column=coursera_details:learner_id, timestamp=1529314684719, value=2
2 row(s) in 0.1430 seconds

hbase(main):005:0> get 'coursera', 'Bigdata'
COLUMN
0 row(s) in 0.1090 seconds

hbase(main):006:0> get 'coursera', '1'
COLUMN
coursera_details:course_duration timestamp=1529314663277, value=6
coursera_details:course_id timestamp=1529314641749, value=123
coursera_details:course_name timestamp=1529314651864, value=Bigdata
coursera_details:learner_id timestamp=1529314629019, value=1
coursera_details:tutor_name timestamp=1529314672629, value=Jack
5 row(s) in 0.0690 seconds

hbase(main):007:0> get 'coursera', '2'
COLUMN
coursera_details:course_duration timestamp=1529314738104, value=7
coursera_details:course_id timestamp=1529314708322, value=124
coursera_details:course_name timestamp=1529315359923, value=Sattistical Learning
coursera_details:learner_id timestamp=1529314684719, value=2
4 row(s) in 0.0520 seconds

hbase(main):008:0>
```

Query 3- DELETE DATA

```

sridevi@sridevi-VirtualBox:~$
2 row(s) in 0.1430 seconds
hbase(main):005:0> get 'coursera', 'Bigdata'
COLUMN
0 row(s) in 0.1090 seconds
hbase(main):006:0> get 'coursera', '1'
COLUMN
coursera_details:course_duration    timestamp=1529314603277, value=6
coursera_details:course_id         timestamp=1529314641749, value=123
coursera_details:course_name       timestamp=1529314651864, value=Bigdata
coursera_details:learner_id        timestamp=1529314629019, value=1
coursera_details:tutor_name        timestamp=1529314672629, value=Jack
5 row(s) in 0.0680 seconds
hbase(main):007:0> get 'coursera', '2'
COLUMN
coursera_details:course_duration    timestamp=1529314738184, value=7
coursera_details:course_id         timestamp=1529314708322, value=124
coursera_details:course_name       timestamp=1529315359923, value=Sattistical Learning
coursera_details:learner_id        timestamp=1529314684719, value=2
4 row(s) in 0.0520 seconds
hbase(main):008:0> delete 'coursera', '1'
ERROR: wrong number of arguments (2 for 3)
Here is some help for this command:
put a delete cell value at specified table/row/column and optionally
timestamp coordinates. Deletes must match the deleted cell's
coordinates exactly. When scanning, a delete cell suppresses older
versions. To delete a cell from 't1' at row 'r1' under column 'c1'
marked with the time 'ts1', do:
hbase> delete 't1', 'r1', 'c1', ts1
hbase(main):009:0> deleteall 'coursera', '1'
0 row(s) in 0.0070 seconds
hbase(main):010:0> scan 'coursera'
COLUMN+CELL
2 row(s) in 0.0890 seconds
2
column=coursera_details:course_duration, timestamp=1529314738184, value=7
2
column=coursera_details:course_id, timestamp=1529314708322, value=124
2
column=coursera_details:course_name, timestamp=1529315359923, value=Sattistical Learning
1 row(s) in 0.0890 seconds
hbase(main):011:0>

```

No Sql Comparison over HBase and Cassandra

The objective behind implementing the No Sql systems is to obtain data volume, and to overcome the limitations of the Relational model.

Cassandra:

- No single point failure
- Scalable, High performance

HBase:

- Open source
- Non-relational

NoSQL Performance can be varied when taken the load and throughput in to consideration.

Operations	CASSANDRA	HBase
Update	Update coursera.coursera_details set course_duartion =5 ,course_name='Introduction to Statstical Learning' where learner_id=3;	put 'coursera', '2', 'coursera_details:course_name', 'Sattistical Learning'
Select	select tutor_name from coursera.coursera_details where course_name='Cloud Computing' ALLOW FILTERING;	get 'coursera', '1' get 'coursera', '2'
Delete	DELETE from	deleteall 'coursera','2'

	courseera.courseera_details where learner_id = 2;	
--	---	--

For the above sample test case data set for coursera when taken in to consideration- **Cassandra's fast write and read performance** has been observed when compared to HBase

Conclusion:

As stated, the above workflow with certain set of parameters is followed in solving the execution by implementing the core and basic concepts of Map reduce algorithm, Cassandra, HBase and Hadoop eco file system.

Source code link:

<https://github.com/PragathiThammaneni/Bigdata-Programming--Hadoop-Spark/tree/master/Lab%201/Source%20Code>

Video Link: <https://youtu.be/cJYrfReYFW0>

Wiki Link:

<https://github.com/PragathiThammaneni/Bigdata-Programming--Hadoop-Spark/wiki/Lab-1-Assignment>