



**UNIVERSITY OF MISSOURI-KANSAS CITY**

**Bigdata Hadoop Programming**

**# Lab 3 Assignment**

**Team Members:**

Pragathi Thammaneni

Sridevi Mallipudi

### **Introduction:**

The main core concept for executing the Lab 3 Assignment is to implement the MapReduce Algorithm for finding Facebook common friends problem and run the MapReduce job on Apache Spark. And also, to implement the Spark Data frames and including intuitive queries like pattern recognition, topic discussion, most important terms etc.

### **Objectives:**

To code for the 2 questions the below concepts are implemented.

Map reduce Algorithm for Facebook common friend's problem

Apache Spark

Spark Data frames

Datasets -queries

### **Approaches /Methods:**

Apache Spark

### **Workflow &Datasets/Parameters and Evaluation:**

The below each question will follow different approaches to solve. Coding is done to perform the evaluation of each individual snippet to execute the datasets which are provided as the input parameters.

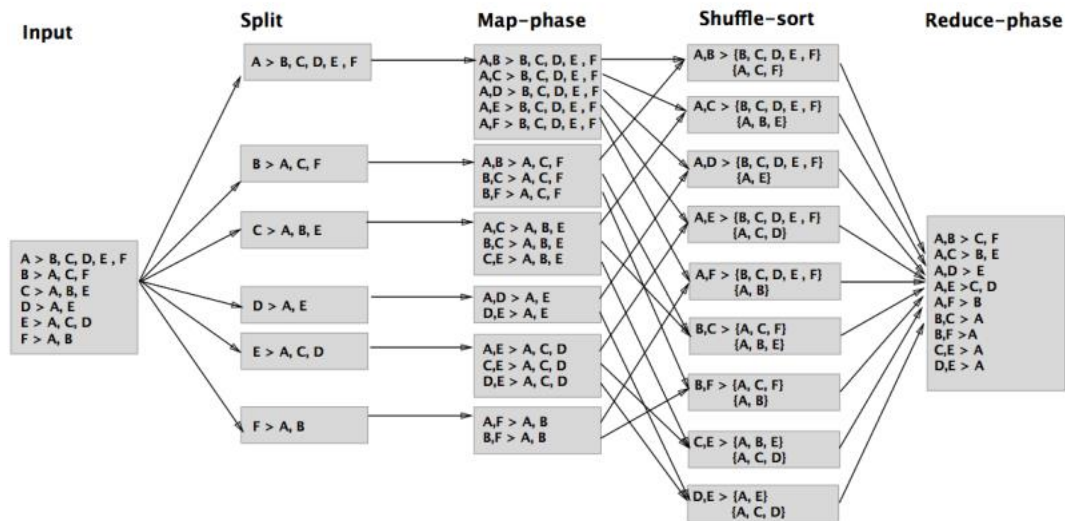
### **Question 1:**

#### **1. Hadoop MapReduce Algorithm**

- a) Implement MapReduce algorithm for finding Facebook common friends problem and run the MapReduce job on Apache Spark.
- b) *Use the following dataset to run your program:*  
<https://umkc.box.com/s/y6juor0fwe96f6louboy3mvbfpli6pgt>
- c) Write a report including your algorithm and result screenshots.

For solving the below question a map reduce algorithm is used to find out the face book common friend problem. Below is the sample use case for implementing the Map reduce.

## Use Case Diagram:



Below is the screen shot for the Map Reduce Algorithm

```

1  import org.apache.spark.sql.SparkSession
2
3  object Friends {
4    def main(args: Array[String]) @Unit {
5
6      System.setProperty("hadoop.home.dir", "F:\\windows\\hadoop")
7
8      val sc = SparkSession
9        .builder
10         .appName("SparkWordCount")
11         .master("local[*]")
12         .getOrCreate().sparkContext
13
14      val flist = sc.textFile(path = "input.txt")
15      val common = flist.flatMap { x =>
16        val splitflist = x.split(regex = " : ")
17        val owner = splitflist(0)
18        val friendslist = splitflist(1).split(regex = " ")
19        friendslist.foreach(println)
20        val makelist = friendslist.slice(0, friendslist.size).map(y => {
21          if (owner > y) (y, owner) else (owner, y)
22        })
23        makelist.map(z => (z, friendslist.slice(0, friendslist.size).toSet))
24      }
25
26      val findcommon = common.reduceByKey(x, y => x intersect y).sortByKey(0)
27      findcommon.collect().take(10).foreach(x => {
28        println(s"${x._1} -> ${x._2.mkString(" ")}")
29        findcommon.saveAsTextFile(path = "output")
30      })
31    }
32  }

```

The screenshot shows the code in the Friends.scala file. The code uses Spark to read an input file, split it into key-value pairs, shuffle and sort by key, and then reduce to find mutual friends. The output is saved to a file named "output".

## Output Screenshot:

```
1  A B C D
2  A C B D
3  A D B C
4  B C A D E
5  B D A C E
6  B E C D
7  C D A B E
8  C E B D
9  D E B C
10
```

## Question 2:

### 2. Spark Data Frames

#### Datasets:

1. **FIFA World Cup:**  
<https://www.kaggle.com/abecklas/fifa-world-cup#WorldCupMatches.csv>
2. **Kickstarter Projects**  
<https://www.kaggle.com/kemical/kickstarter-projects>
3. **Google-Landmarks Dataset**  
<https://www.kaggle.com/google/google-landmarks-dataset>

- a. Create a Spark DataFrame using one of datasets, trying to use all different StructType.
- b. Perform 10 intuitive questions in Dataset (e.g.: pattern recognition, topic discussion, most important terms, etc.). Use your innovation to think out of box.
- c. Perform any 5 queries in Spark RDD's and Spark Data Frames. Compare the results

#### Queries implemented:

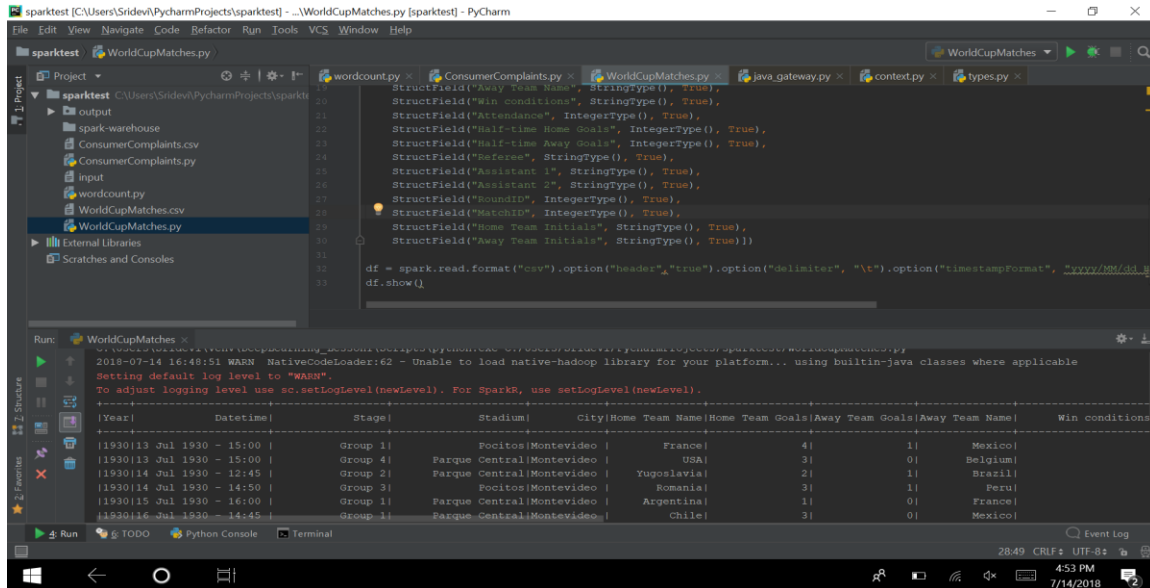
Below are the queries that has been implemented on the data sets FIFA World Cup.

1. Creation of Spark Data Frame using FIFA World Cup.
2. Queries regarding pattern recognition, topic discussion and most important terms has been performed.

3. Performed queries using Spark RDD's and Spark Data Frames.

Please refer the source code folder for the queries description

## Struct Type:



```
sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatches.py [sparktest] - PyCharm
File Edit View Navigate Code Refactor Run Tools VCS Window Help

Project: sparktest
  sparktest
  output
  spark-warehouse
  ConsumerComplaints.csv
  ConsumerComplaints.py
  input
  wordcount.py
  WorldCupMatches.csv
  WorldCupMatches.py
  External Libraries
  Scratches and Consoles

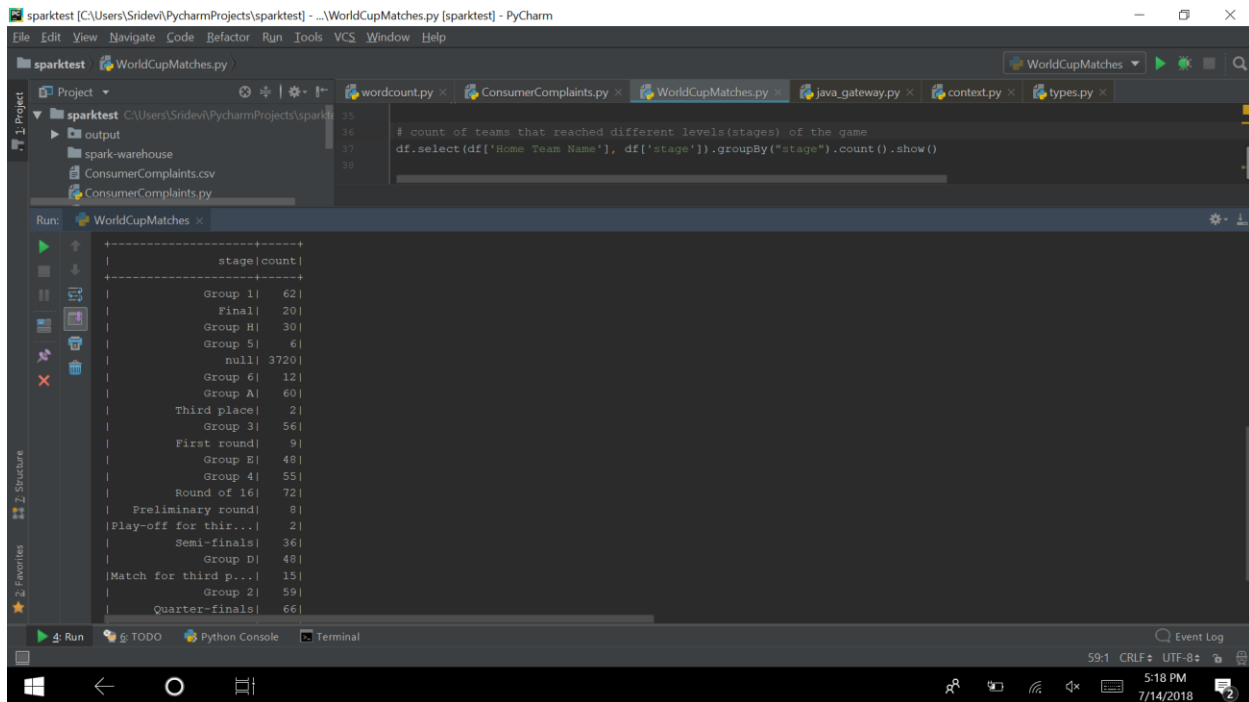
WorldCupMatches.py
20 StructField("Away Team Name", StringType(), True),
21 StructField("Win conditions", StringType(), True),
22 StructField("Attendance", IntegerType(), True),
23 StructField("Half-time Home Goals", IntegerType(), True),
24 StructField("Half-time Away Goals", IntegerType(), True),
25 StructField("Referee", StringType(), True),
26 StructField("Assistant 1", StringType(), True),
27 StructField("Assistant 2", StringType(), True),
28 StructField("RoundID", IntegerType(), True),
29 StructField("MatchID", IntegerType(), True),
30 StructField("Home Team Initials", StringType(), True),
31 StructField("Away Team Initials", StringType(), True))
32
33 df = spark.read.format("csv").option("header", "true").option("delimiter", "\t").option("timestampFormat", "yyyy/MM/dd")
34 df.show()

Run: WorldCupMatches
2018-07-14 16:48:51 WARN NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Year|      Datetime|      Stage|      Stadium|      City|Home Team Name|Home Team Goals|Away Team Goals|Away Team Name|      Win conditions|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|1930|13 Jul 1930 - 15:00|      Group 1|      Pocitos|Montevideo|      France|          4|          1|      Mexico|
|1930|13 Jul 1930 - 15:00|      Group 4|      Parque Central|Montevideo|      USA|          3|          0|      Belgium|
|1930|14 Jul 1930 - 12:45|      Group 2|      Parque Central|Montevideo|      Yugoslavia|          2|          1|      Brazil|
|1930|14 Jul 1930 - 14:50|      Group 3|      Pocitos|Montevideo|      Romania|          3|          1|      Peru|
|1930|15 Jul 1930 - 16:00|      Group 1|      Parque Central|Montevideo|      Argentina|          1|          0|      France|
|1930|16 Jul 1930 - 14:45|      Group 1|      Parque Central|Montevideo|      Chile|          3|          0|      Mexico|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

28:49 CRLF UTF-8 4:53 PM 7/14/2018
```

## 10 intuitive queries:

### Query 1:Count of teams that reached different levels(stages) of the game



```
sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatches.py [sparktest] - PyCharm
File Edit View Navigate Code Refactor Run Tools VCS Window Help

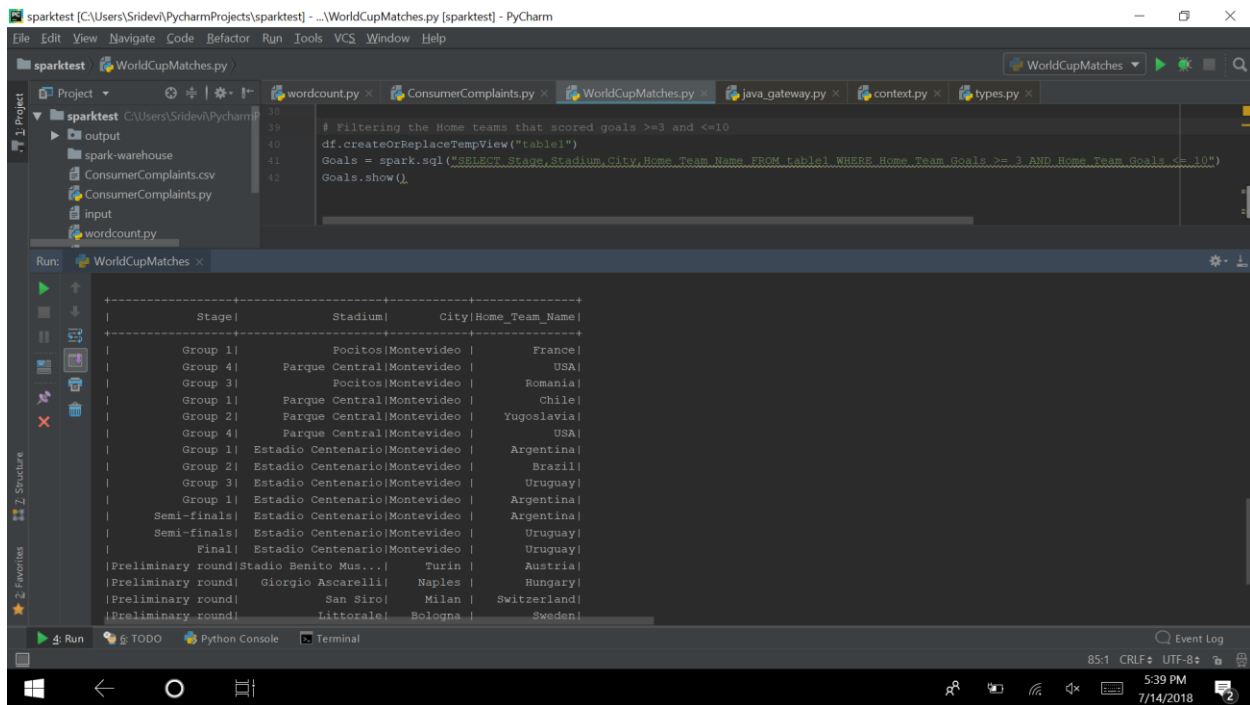
Project: sparktest
  sparktest
  output
  spark-warehouse
  ConsumerComplaints.csv
  ConsumerComplaints.py
  input
  wordcount.py
  WorldCupMatches.csv
  WorldCupMatches.py
  External Libraries
  Scratches and Consoles

WorldCupMatches.py
35
36 # count of teams that reached different levels(stages) of the game
37 df.select(df['Home Team Name'], df['stage']).groupBy("stage").count().show()
38

Run: WorldCupMatches
+-----+-----+
|stage|count|
+-----+-----+
|Group 1| 62|
|Final| 20|
|Group H| 30|
|Group 5| 6|
|null| 3720|
|Group 6| 12|
|Group A| 60|
|Third place| 2|
|Group 3| 56|
|First round| 9|
|Group E| 48|
|Group 4| 55|
|Round of 16| 72|
|Preliminary round| 8|
|Play-off for third place| 2|
|Semi-finals| 36|
|Group D| 48|
|Match for third place| 15|
|Group 2| 59|
|Quarter-finals| 66|
+-----+-----+

59:1 CRLF UTF-8 5:18 PM 7/14/2018
```

## Query 2: Filtering the Home teams that scored goals $\geq 3$ and $\leq 10$

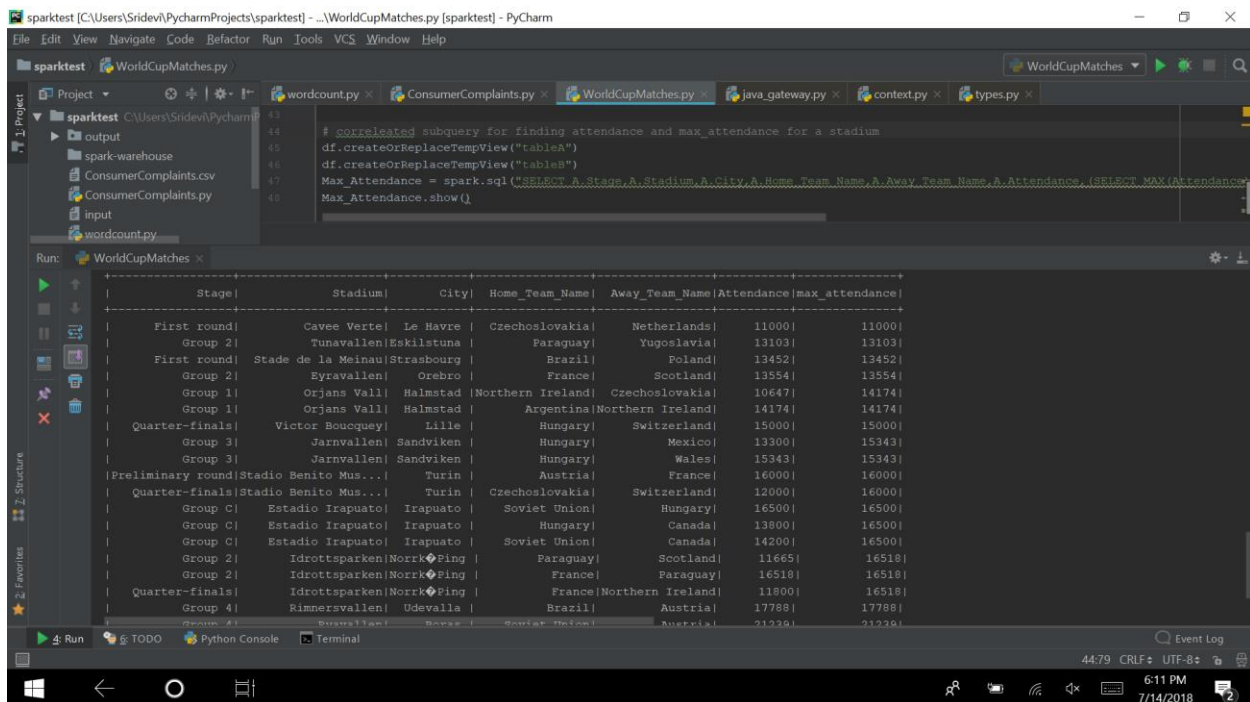


The screenshot shows the PyCharm IDE with a project named 'sparktest'. The file 'WorldCupMatches.py' is open, containing a Spark SQL query. The query filters home teams based on the number of goals scored, specifically those with 3 or more goals and 10 or fewer goals. The output of the query is displayed in a table format within the IDE's run console.

```
38 # Filtering the Home teams that scored goals >=3 and <=10
39 df.createOrReplaceTempView("table1")
40 Goals = spark.sql("SELECT Stage, Stadium, City, Home_Team_Name FROM table1 WHERE Home_Team_Goals >= 3 AND Home_Team_Goals <= 10")
41 Goals.show()
```

Stage	Stadium	City	Home_Team_Name
Group 1	Pocitos	Montevideo	France
Group 4	Parque Central	Montevideo	USA
Group 3	Pocitos	Montevideo	Romania
Group 1	Parque Central	Montevideo	Chile
Group 2	Parque Central	Montevideo	Yugoslavia
Group 4	Parque Central	Montevideo	USA
Group 1	Estadio Centenario	Montevideo	Argentina
Group 2	Estadio Centenario	Montevideo	Brazil
Group 3	Estadio Centenario	Montevideo	Uruguay
Group 1	Estadio Centenario	Montevideo	Argentina
Semi-finals	Estadio Centenario	Montevideo	Argentina
Semi-finals	Estadio Centenario	Montevideo	Uruguay
Final	Estadio Centenario	Montevideo	Uruguay
Preliminary round	Stadio Benito Mus...	Turin	Austria
Preliminary round	Giorgio Ascarelli	Naples	Hungary
Preliminary round	San Siro	Milan	Switzerland
Preliminary round	Littorale	Bologna	Sweden

## Query 3: Correlated subquery for finding attendance and max\_attendance for a stadium

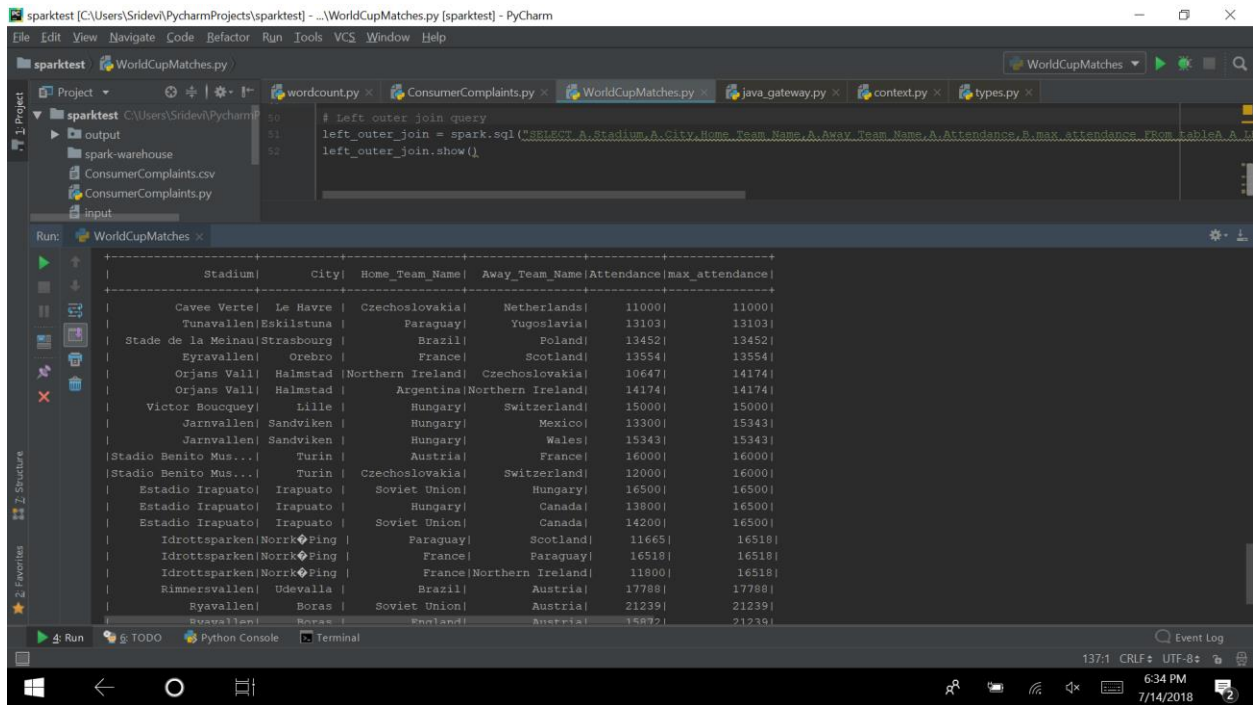


The screenshot shows the PyCharm IDE with a project named 'sparktest'. The file 'WorldCupMatches.py' is open, containing a Spark SQL query. The query uses a correlated subquery to find the attendance and maximum attendance for each stadium. The output of the query is displayed in a table format within the IDE's run console.

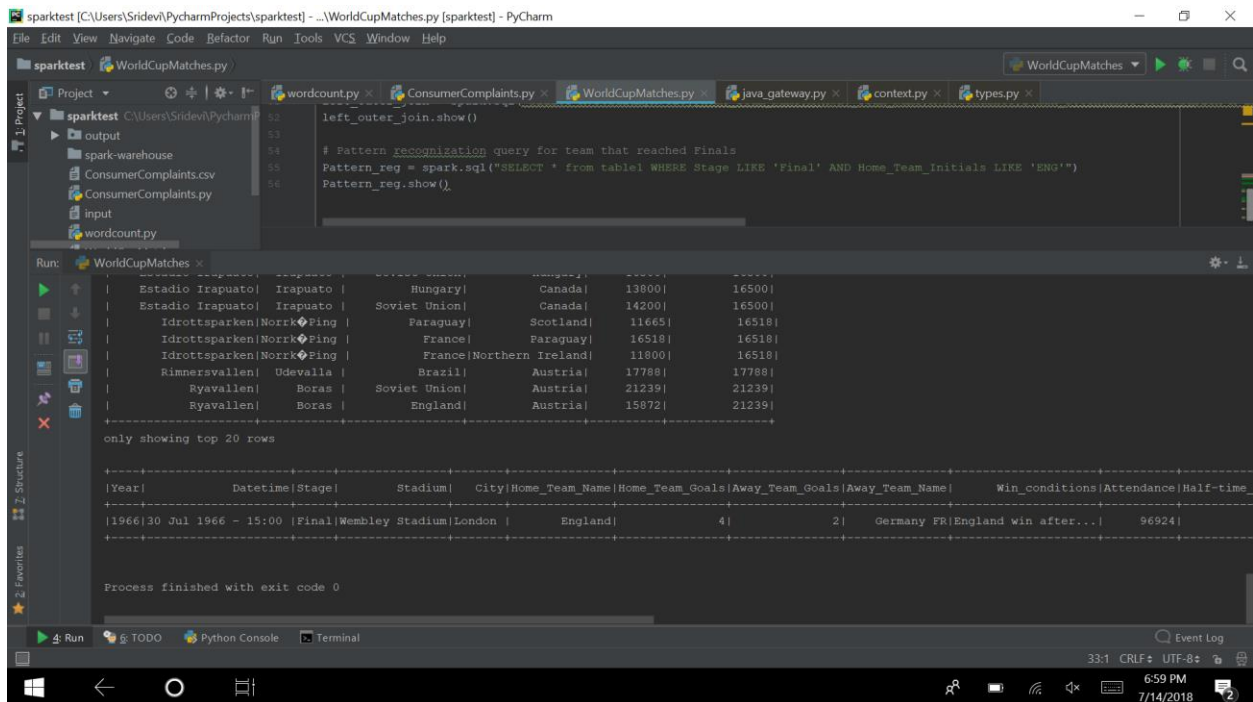
```
43 # correlated subquery for finding attendance and max_attendance for a stadium
44 df.createOrReplaceTempView("tableA")
45 df.createOrReplaceTempView("tableB")
46 Max_Attendance = spark.sql("SELECT A.Stage, A.Stadium, A.City, A.Home_Team_Name, A.Away_Team_Name, A.Attendance, (SELECT MAX(Attendance) FROM tableA WHERE Stadium = A.Stadium) AS Max_Attendance FROM tableA")
47 Max_Attendance.show()
```

Stage	Stadium	City	Home_Team_Name	Away_Team_Name	Attendance	max_attendance
First round	Cavee Verte	Le Havre	Czechoslovakia	Netherlands	11000	11000
Group 2	Tunavallen	Eskilstuna	Paraguay	Yugoslavia	13103	13103
First round	Stade de la Meinau	Strasbourg	Brazil	Poland	13452	13452
Group 2	Eyravallen	Orebro	France	Scotland	13554	13554
Group 1	Orjens Vall	Halmstad	Northern Ireland	Czechoslovakia	10647	14174
Group 1	Orjens Vall	Halmstad	Argentina	Northern Ireland	14174	14174
Quarter-finals	Victor Boucouey	Lille	Hungary	Switzerland	15000	15000
Group 3	Jarnvallen	Sandviken	Hungary	Mexico	13300	15343
Group 3	Jarnvallen	Sandviken	Hungary	Wales	15343	15343
Preliminary round	Stadio Benito Mus...	Turin	Austria	France	16000	16000
Quarter-finals	Stadio Benito Mus...	Turin	Czechoslovakia	Switzerland	12000	16000
Group C	Estadio Irapuato	Irapuato	Soviet Union	Hungary	16500	16500
Group C	Estadio Irapuato	Irapuato	Hungary	Canada	13800	16500
Group C	Estadio Irapuato	Irapuato	Soviet Union	Canada	14200	16500
Group 2	Idrottsparken	Norrkoping	Paraguay	Scotland	11665	16518
Group 2	Idrottsparken	Norrkoping	France	Paraguay	16518	16518
Quarter-finals	Idrottsparken	Norrkoping	France	Northern Ireland	11800	16518
Group 4	Rimmersvallen	Udevalla	Brazil	Austria	17788	17788
Group 4	Rimmersvallen	Udevalla	Sweden	Austria	21236	21236

## Query 4: Left outer join query



### Query 5: Pattern recognition query for team that reached Finals



### Query 6: Average goals scored by a team



The screenshot shows a PyCharm IDE with a project named 'sparktest'. The file 'WorldCupMatches.py' is open, containing a Spark SQL query to calculate the average goals scored by each team. The query is: `Avg_goals = spark.sql("SELECT Home_Team_Name AS Team, ROUND(AVG(Home_Team_Goals),0) AS average_goals FROM table1 GROUP BY Home_Team_Name")`. The output of the query is displayed in the Run console, showing a list of teams and their average goals.

Team	average_goals
Paraguay	1.0
Russia	3.0
Senegal	2.0
Sweden	2.0
IR Iran	0.0
Turkey	5.0
Salvador	0.0
Iraq	1.0
Germany	2.0
France	2.0
Greece	1.0
Algeria	1.0
Togo	0.0
Slovakia	2.0
Argentina	2.0
Wales	2.0
Belgium	2.0

## Query 7: Count of rows after performing UNION ALL operation

The screenshot shows the same PyCharm IDE with the 'WorldCupMatches.py' file. The code now includes a query to count the total rows after a UNION ALL operation between two tables. The query is: `Union_all = spark.sql("SELECT COUNT(*) AS total_rows FROM (SELECT * FROM tableA UNION ALL SELECT * FROM tableB)sub")`. The output in the Run console shows the total count of rows.

total_rows
1704

Process finished with exit code 0

## Query 8: Number of times a country scores goals greater than or equal to 4



sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatches.py [sparktest] - PyCharm

```
File Edit View Navigate Code Refactor Run Tools VCS Window Help
```

WorldCupMatches.py

```
68 join_all.show()
69
69 # number of times a country scores goals greater than or equal to 4
70 df.createOrReplaceTempView("table3")
71 query = spark.sql("SELECT Home_Team_Name AS COUNTRY,COUNT(Home_Team_Goals) AS No_of_Times FROM table3 where Home_Team_Goals >=4 OR")
72 query.show()
```

Run: WorldCupMatches

COUNTRY	No_of_Times
Brazil	18
Germany FR	11
Hungary	10
Argentina	9
France	8
Germany	6
Uruguay	6
Yugoslavia	5
Italy	5
Portugal	3
Netherlands	3
Soviet Union	3
England	2
Spain	2
Czechoslovakia	2
Switzerland	2
Austria	2
Sweden	2

PEP 8: line too long (169 > 120 characters)

71:168 CRLF UTF-8 8:17 PM 7/14/2018

## Query 9: Number of maximum number of goals scored by a country between years 2000 nd 2010

sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatches.py [sparktest] - PyCharm

```
File Edit View Navigate Code Refactor Run Tools VCS Window Help
```

WorldCupMatches.py

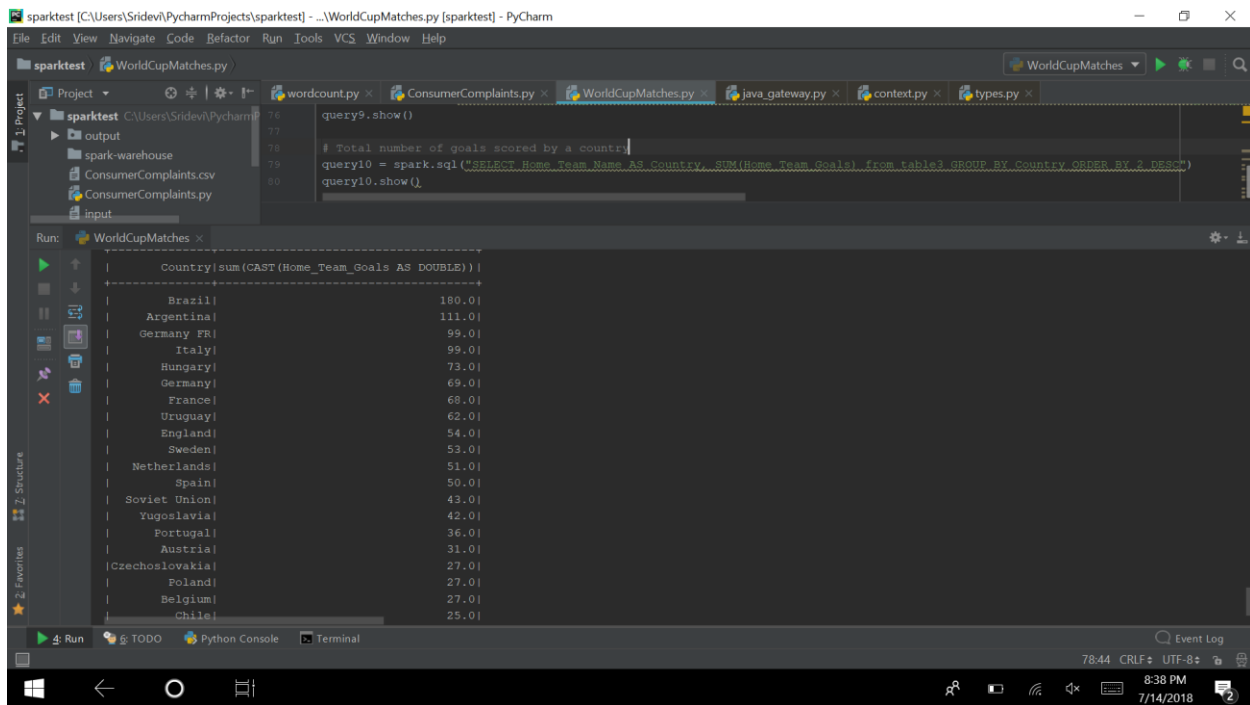
```
73 query9.show()
74
74 # number of maximum number of goals scored by a country between years 2000 nd 2010
75 query9 = spark.sql("SELECT Home_Team_Name, MAX(Home_Team_Goals) from table3 where Year BETWEEN 2000 AND 2010 GROUP BY Home_Team_Name")
76 query9.show()
```

Run: WorldCupMatches

Home_Team_Name	max(Home_Team_Goals)
Germany	8
Portugal	7
Argentina	6
Brazil	4
Spain	4
Italy	3
Cote d'Ivoire	3
Belgium	3
Mexico	3
Poland	3
Ecuador	3
Turkey	3
Slovakia	3
USA	3
Senegal	3
Australia	3
Nigeria	2
Denmark	2
Sweden	2

74:83 CRLF UTF-8 8:32 PM 7/14/2018

## Query 10: Total number of goals scored by a country



```
sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatches.py [sparktest] - PyCharm
File Edit View Navigate Code Refactor Run Tools VCS Window Help

sparktest WorldCupMatches.py
Project sparktest C:\Users\Sridevi\PycharmProjects\sparktest
output
spark-warehouse
ConsumerComplaints.csv
ConsumerComplaints.py
input

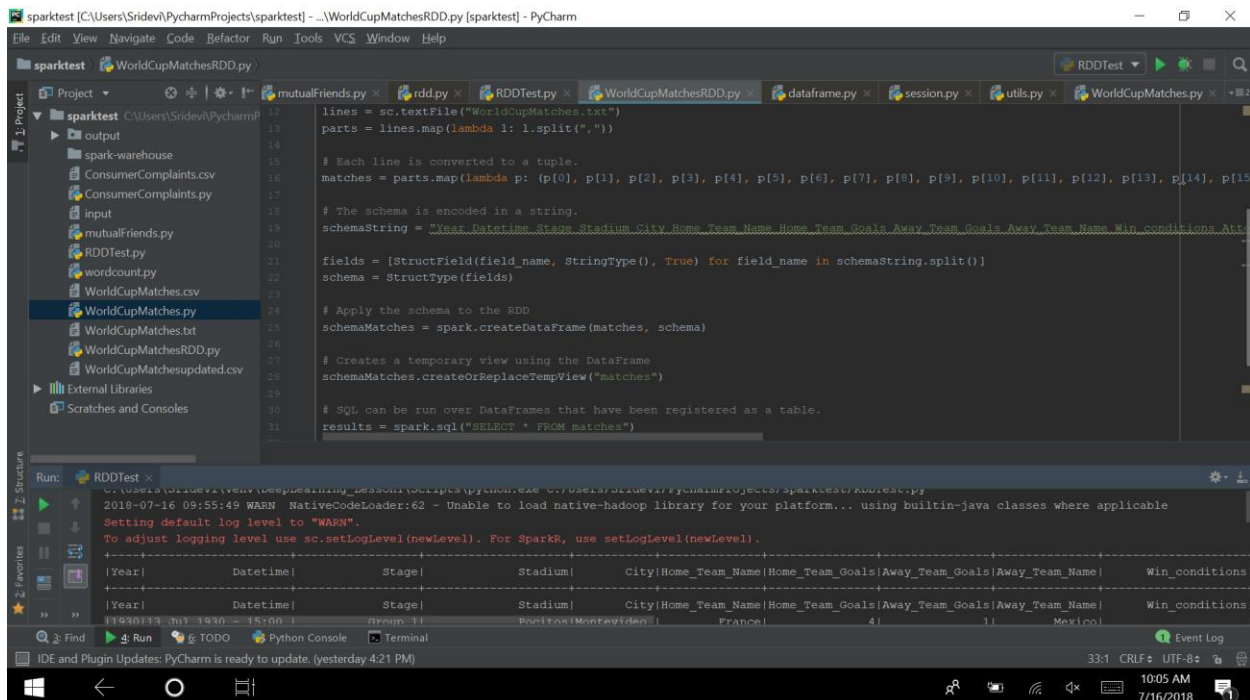
76 query9.show()
77
78 # Total number of goals scored by a country
79 query10 = spark.sql("SELECT Home_Team_Name AS Country, SUM(Home_Team_Goals) from table3 GROUP BY Country ORDER BY 2 DESC")
80 query10.show()
```

Run: WorldCupMatches x

Country	sum(CAST(Home_Team_Goals AS DOUBLE))
Brazil	180.0
Argentina	111.0
Germany FR	99.0
Italy	99.0
Hungary	73.0
Germany	69.0
France	68.0
Uruguay	62.0
England	54.0
Sweden	53.0
Netherlands	51.0
Spain	50.0
Soviet Union	43.0
Yugoslavia	42.0
Portugal	36.0
Austria	31.0
Czechoslovakia	27.0
Poland	27.0
Belgium	27.0
Chile	25.0

78:44 CRLF UTF-8 8:38 PM 7/14/2018

## Spark RDD's: Apply the schema to the RDD



```
sparktest [C:\Users\Sridevi\PycharmProjects\sparktest] - ...WorldCupMatchesRDD.py [sparktest] - PyCharm
File Edit View Navigate Code Refactor Run Tools VCS Window Help

sparktest WorldCupMatchesRDD.py
Project sparktest C:\Users\Sridevi\PycharmProjects\sparktest
output
spark-warehouse
ConsumerComplaints.csv
ConsumerComplaints.py
input
mutualFriends.py
RDDTest.py
wordcount.py
WorldCupMatches.csv
WorldCupMatches.py
WorldCupMatches.txt
WorldCupMatchesRDD.py
WorldCupMatchesupdated.csv
External Libraries
Scratches and Consoles

12 lines = sc.textFile("WorldCupMatches.txt")
13 parts = lines.map(lambda l: l.split(","))
14
15 # Each line is converted to a tuple.
16 matches = parts.map(lambda p: (p[0], p[1], p[2], p[3], p[4], p[5], p[6], p[7], p[8], p[9], p[10], p[11], p[12], p[13], p[14], p[15])
17
18 # The schema is encoded in a string.
19 schemaString = "Year,Datetime,Stage,Stadium,City,Home_Team_Name,Home_Team_Goals,Away_Team_Goals,Away_Team_Name,Win_conditions,Att"
20
21 fields = [StructField(field_name, StringType(), True) for field_name in schemaString.split(",")]
22 schema = StructType(fields)
23
24 # Apply the schema to the RDD
25 schemaMatches = spark.createDataFrame(matches, schema)
26
27 # Creates a temporary view using the DataFrame
28 schemaMatches.createOrReplaceTempView("matches")
29
30 # SQL can be run over DataFrames that have been registered as a table.
31 results = spark.sql("SELECT * FROM matches")
32
```

Run: RDDTest x

2018-07-16 09:55:49 WARN NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

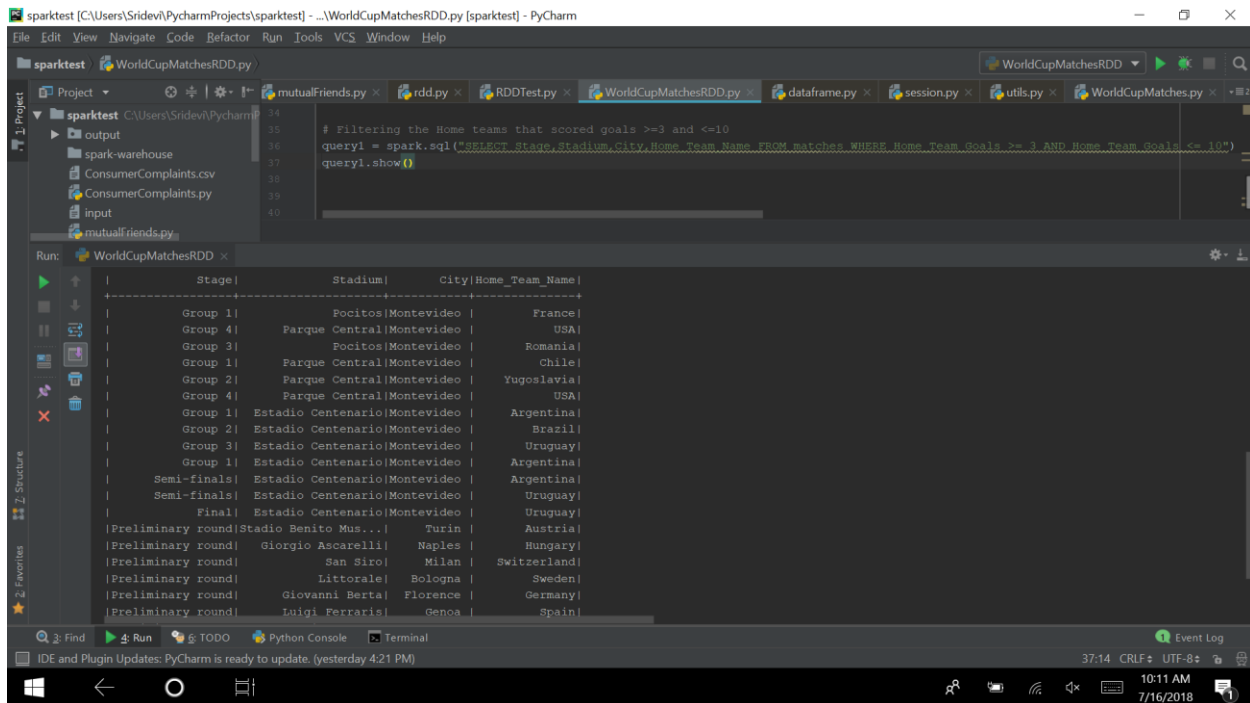
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

Year	Datetime	Stage	Stadium	City	Home_Team_Name	Home_Team_Goals	Away_Team_Goals	Away_Team_Name	Win_conditions
1930	1930-07-13	Final	Estadio Centenario	Montevideo	Uruguay	4	2	Uruguay	W

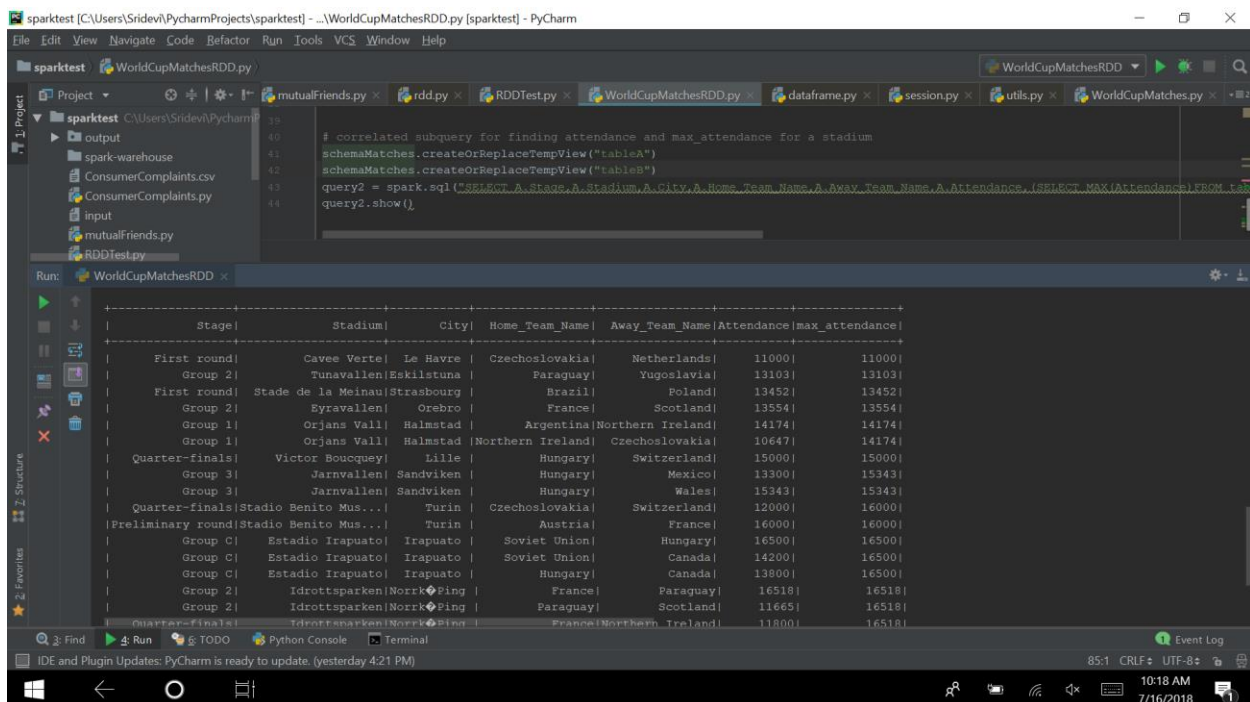
IDE and Plugin Updates: PyCharm is ready to update. (yesterday 4:21 PM)

33:1 CRLF UTF-8 10:05 AM 7/16/2018

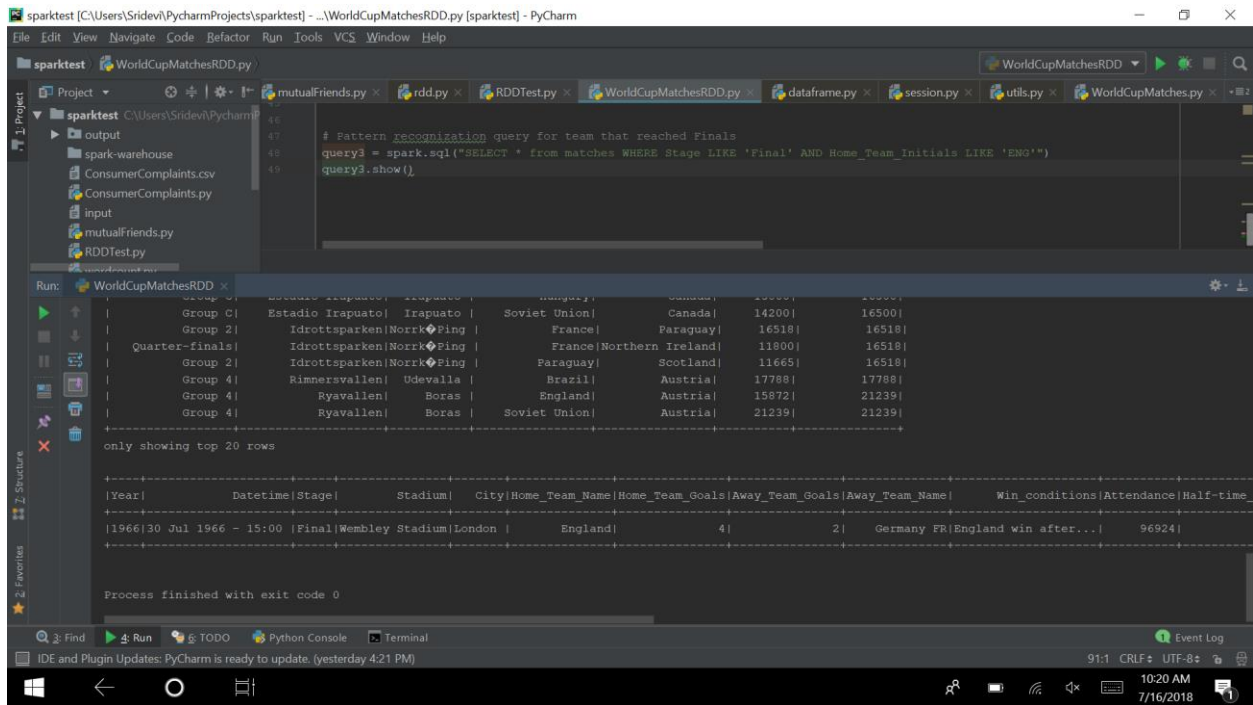
## Query 1: Filtering the Home teams that scored goals >=3 and <=10 with RDD



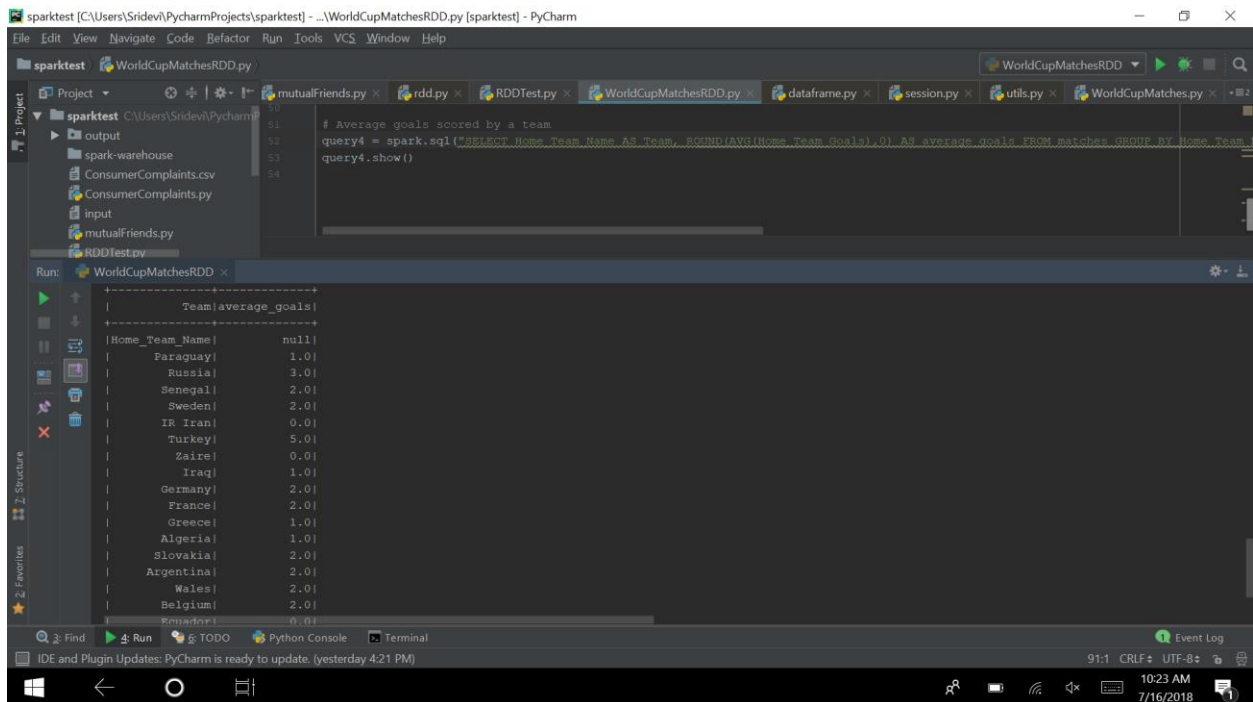
## Query 2: Correlated subquery for finding attendance and max\_attendance for a stadium with RDD



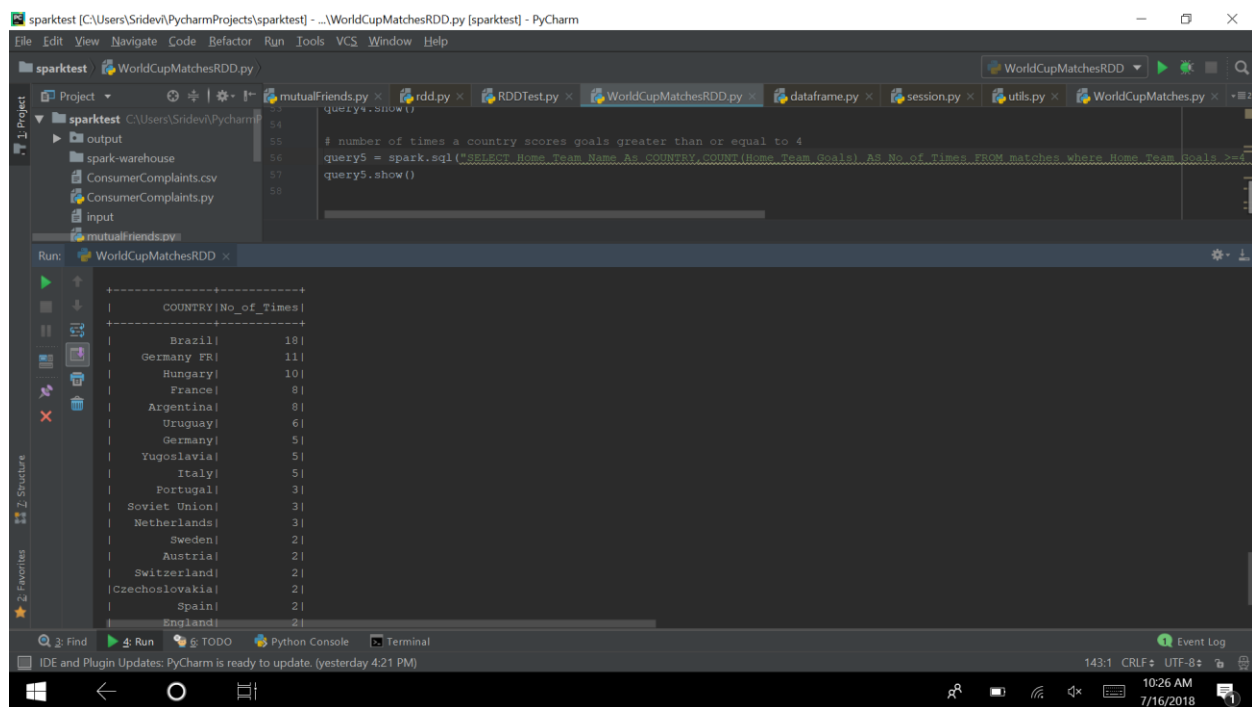
## Query 3: Pattern recognition query for team that reached Finals with RDD



#### Query 4: Average goals scored by a team with RDD



#### Query 5: Number of times a country scores goals greater than or equal to 4 with RDD



### Comparison between Spark RDD's and the data frames:

Spark RDD's	Spark Data Frames
RDD is a fault-tolerant collection of elements that can be operated on in-parallel, also we can say RDD is the fundamental data structure of Spark	It is a distributed collection of data. Basically, data is organized into named columns in data frames.
Basically, it is read-only partition collection of records. Moreover, it supports in-memory computations on large clusters in a fault-tolerant manner	Spark also introduced <b>catalyst optimizer</b> , along with data frame. To build an extensible query optimizer, it also leverages advanced programming features.
An RDD can come easily handle data with no predefined structure.	In Spark, data frame allows developers to impose a structure onto a distributed data. It also allows higher-level abstraction.
Compile- Time Type Safety RDD- RDD Supports <i>object-oriented programming</i> style with compile-time type safety.	If we try to access any column which is not present in the table, then an attribute error may occur at runtime. Data frame will not support compile-time type safety in such case.
Spark does not compute their result right away, it evaluates RDDs lazily	Computation happens only when action appears as Spark evaluates data frame lazily

RDDs of Apache spark offers low-level functionality and control. Whereas datasets offer higher functionality. While data frame offers high-level domain-specific operations, saves space and executes at high speed. Therefore, it increases the efficiency of the system.

**Conclusion:**

As stated, the above workflow with certain set of parameters is followed in solving the execution by implementing the core and basic concepts Apache Spark.

**Source code link:**

<https://github.com/PragathiThammaneni/Bigdata-Programming--Hadoop-Spark/tree/master/Lab%203>

**Video Link: Provided in wiki** <https://youtu.be/31C5DO0zQSE>

**Wiki Link:**

<https://github.com/PragathiThammaneni/Bigdata-Programming--Hadoop-Spark/wiki/Lab-3-Assignment>